

What a Transformer Retrieves and What It Computes: A Measured Theory of the Composition Core, with Incidence Calculus, Tropical Geometry, and Semiring Datalog as Three Interpretations

J. Allan Scott
Independent Researcher
allan@jallan.scot

[Second Author — to be confirmed]
[Affiliation — seeking an established co-author]

Abstract

We decompile small transformer language models into a flat *retrieval* store (n-gram, induction, grammar) plus a *composition* kernel, and use the decomposition to ask a sharp, measurable question of every next-token decision: is the model *retrieving* a symbolic rule, *selecting* within a rule-proposed set, or *computing* something new? Across two Qwen2.5-0.5B models and the four-rung Pythia ladder (70M–1B), the labour splits roughly 25% retrieved, 60% selected, 15% composed: a small but irreducible core of genuinely computed tokens, the *irreducible computation*, that grows with scale. The mechanism is not selection by magnitude: the predicted logit is a uniformly-distributed ~ 45 -way additive sum (participation ratio $PR \approx 45$, route-invariant). Retrieval versus computation instead separates cleanly on two near-orthogonal, exactly measurable axes: a *normalised power-diagram margin* (the Euclidean distance from the residual to the nearest unembedding facet) and a *single-circuit readout multiplicity* μ_r . Composed tokens are *emergent* (the argmax of the sum but of no summand) and *causally fragile* under ablation, with a repair term that is provably diffuse ($P(\text{single-module rescue}) \approx 1/PR$). We then give one theory in three interpretations. It accounts for these measurements and turns out to be a single object at two temperatures. As a probabilistic *logic* (temperature one) it is Bundy’s incidence calculus with two assumptions removed: incidences become signed measures in an inner-product space and Boolean $\cap/\cup$ becomes a weighted threshold, with the Gram kernel $G_{vw} = \langle U_v, U_w \rangle$ as the explicit carrier of non-truth-functionality. As a *geometry* (temperature zero) it is the tropical variety of the decision surface, and the irreducible computation is a tropical-rank floor that linear rank cannot see. As a *computation* it is a semiring-weighted Datalog program whose two temperatures are sum-product (the softmax) and max-product (the greedy decode) over one program; Maslov dequantization is the homomorphism between them. The retrievable fragment exports to compact, verifiable Datalog; the irreducible computation is exactly the dense-Gram, high-treewidth, high-tropical-rank region that no lookup table reproduces. Every quantitative claim is reproducible with the

fieldrun toolchain, to be released with the camera-ready version.

1. Introduction

A large language model, asked for the next token, sometimes does something a lookup table could do and sometimes does not. “The capital of France is” is answered by retrieval; the model has effectively memorised a fact and the harder question is only where it is stored [11, 19]. But many tokens are not in any table the model could be said to consult: they are *computed* on the fly, assembled from a distributed interaction of partial contributions. The central empirical question of this paper is how to tell these apart *operationally*, on a real model, token by token, and what the mechanism of each looks like.

We attack the question with a *decompiler*. The fieldrun toolchain splits a small transformer into two halves that together reproduce it: a flat *retrieval* store (n-gram successor tables, induction copy-rules, a closed-class grammar skeleton) that reconstructs roughly half of the model’s predictions with no linear algebra at all, and a *composition* kernel that runs the attention+MLP forward pass as plain matrix multiplies over flat weight arrays [23]. Every weight tier is validated by exact top-1 agreement against a reference implementation, so the decomposition is faithful rather than approximate. Given the retrieval store’s candidate set for a context, we classify the model’s argmax token t into three routes: RETRIEVED (a single store idiom’s top-1 is t , a pure symbolic lookup), SELECTED (t is in the candidate set but is no idiom’s top-1; the set contains the answer, the choice within it is made elsewhere), and COMPOSED (t is in no rule’s output, the *irreducible computation*, genuinely computed).

The headline measurement is stable and, we will argue, theoretically loaded. Across two Qwen2.5-0.5B models [24] and all four rungs of the Pythia ladder [2], the model’s labour divides roughly 25% RETRIEVED, 60% SELECTED, 15% COMPOSED. Composition is mostly *disambiguation within a retrieved set*, not generation from nothing; only the $\sim 15\%$ irreducible computation is from scratch, and it grows monoton-

ically with model scale against a fixed store. Three further facts sharpen the picture and resist the obvious explanations. First, the route split is *not* a magnitude distinction: the predicted logit is a uniformly distributed ~ 45 -way additive sum (participation ratio $PR \approx 45$, route-invariant), so no circuit ever “decides” a token by dominating the sum, not even one a single n -gram rule reproduces perfectly. Second, the split separates cleanly on two other, near-orthogonal axes that we can measure *exactly*: a decision *margin* that is the Euclidean distance from the residual stream to the nearest unembedding power-diagram facet, and a single-circuit *readout multiplicity* μ_t . Third, composed tokens are *emergent* (the argmax of the sum but of no summand, $\mu_t \approx 0$) and *causally fragile* under ablation, with a repair mechanism that is provably diffuse rather than localisable.

These measurements ask for a theory, and the theory we give is one object seen from three sides. The thread that ties them is older than transformers: it is Alan Bundy’s *incidence calculus* [4], invented in 1985 precisely to reason with probabilities that are *not truth-functional* by tracking the underlying *incidences* (the worlds in which a proposition holds) and recovering probability as a measure on them. We make three claims, each an interpretation of the same structure:

1. As a probabilistic **logic** at temperature $T = 1$ (the soft accumulation and the recovered output measure), the core is incidence calculus with exactly two of Bundy’s assumptions removed. Incidences become *signed measures in an inner-product space* rather than sets, and Boolean $\cap/\cup$ become a *weighted threshold*. The Gram kernel $G_{vw} = \langle U_v, U_w \rangle$ of the unembedding frame is the explicit operator carrying non-truth-functionality; the classical Boolean calculus is recovered *exactly* as the diagonal- G limit.
2. As a **geometry** at temperature $T = 0$ (the hard argmax, its cells, margins and rank), the decision surface is the *tropical variety* of a max-logit polynomial whose monomials are the unembedding rows. Composition is interior (non-monomial) tropical points, and the irreducible computation is a *tropical-rank floor*: the gap between the model’s tropical rank and that of any flat lookup table, a gap that linear (SVD) rank structurally cannot measure.
3. As a **computation** (the executable, statically checkable artifact), the core is a *semiring-weighted Datalog program* Π . Evaluated under the log-semiring it is sum-product and returns the softmax distribution; under the tropical semiring it is max-product and returns the greedy decode. The two temperatures are one program under two semirings, and *Maslov dequantization* is the homomorphism between them.

The unification is not loose analogy. The same kernel G that hardens near-synonym competition in the logic is the facet geometry in the tropical picture and the provenance coupling in the Datalog picture; the same irreducible computation is the high participation-ratio residual, the high tropical rank, and

the high-treewidth dense- G region; the same temperature parameter is the semiring choice. We present the measurements first (Sections 3–4), because they constrain every axiom of the theory, and then the three interpretations and their unification (Section 5), the reproduction recipe (Section 6), and the open frontier (Section 7). The tribute is plain throughout: this is Bundy’s calculus, forty years on, meeting a substrate he could not have had.

2. Background and Related Work

2.1 Mechanistic interpretability

Our measurement apparatus inherits the circuits program’s central move: read computation off the model in the basis of its own weights. The residual stream is a linear sum of component writes, so a final-layer logit decomposes additively over heads and neurons; the *direct logit attribution* (DLA) of a component to token v is $\langle d_j, U_v \rangle$, the inner product of the component’s write d_j with the unembedding row U_v [10]. *Superposition*, the storage of more features than dimensions as a near-orthogonal overcomplete code, predicts exactly the distributed, low-per-circuit-magnitude readout we measure [9], and sparse autoencoders attempt to recover that code as monosemantic features [3, 7]. *Induction heads*, the copy-from-earlier-match mechanism behind much in-context learning [20], are the one idiom that exports as a clean recursive rule (Section 5). Feed-forward layers as key–value memories [11] and the locate-and-edit line [19] are the retrieval end of the spectrum we are trying to delimit from the computed end. Where this work differs is in making the retrieve-versus-compute distinction a *measured, per-token classification* with an exact geometric and causal characterisation, rather than a case study of one circuit.

2.2 Incidence calculus

Incidence calculus [4, 5] was built on one load-bearing observation: probability is *not* truth-functional, since $P(A \wedge B)$ is not a function of $P(A)$ and $P(B)$, but *incidence* is. Bundy attaches to each proposition A a set $i(A) \subseteq I$ of “incidences” (possible worlds, or sampled situations) in which A holds, makes the connectives set operations ($i(A \wedge B) = i(A) \cap i(B)$), and recovers probability as the measure $P(A) = |i(A)|/|I|$. The joint that probability loses is carried explicitly by the set overlap. Our claim is that the transformer’s composition core is this calculus with the *orthogonality assumption removed*: its propositions are not Boolean indicator sets but directions in an inner-product space whose pairwise overlaps are a dense Gram kernel, and its connective is not intersection but a weighted threshold. The classical calculus reappears, exactly, when that kernel is diagonal.

2.3 Log-linear, tropical, and semiring views

The additive core is a log-linear / product-of-experts model [13, 15], and Markov logic networks [21] are its closest symbolic-probabilistic cousin; what we add is the inner-product (frame-Gram) coupling and the explicit retrievable/computed split.

The hard decision surface is tropical: an argmax over linear forms is the $(\max, +)$ semiring, and the ReLU/PWL-network-to-tropical lineage [25, 17] gives the decision map as a tropical rational function whose linear-region count is bounded by Newton polytope vertices. Tropical *rank* [8] supplies the irreducible-computation floor; power (Laguerre) diagrams [1] are the decision cells. The bridge between the soft ($T=1$) and hard ($T=0$) pictures is *Maslov dequantization*, the $T \rightarrow 0$ limit of the log-semiring into $(\max, +)$ [16]. The executable form lives in provenance / semiring Datalog [12, 14, 6]: sum-product over a semiring is database evaluation, and the temperature is the semiring parameter. The correlated-alternatives regime connects to discrete-choice econometrics (GEV, nested logit) [18]. Finally, the phenomenon of “redundant distributed voting shading to emergent combination” has ingredients with precedent (distributed connectionist production systems [22], superposition, products of experts, and incidence calculus over a continuous learned space), but, as far as we are aware, the fusion is new.

3. The Decompiler and the Measurement

3.1 The two-tier decompile

`fieldrun` is a single static binary that runs a model from a flat *bundle* (a raw weight blob plus a small manifest) with no deep-learning framework at runtime. The build side decompiles a model into **Tier A**, a retrieval store (induction thresholds; quad/tri/bi/uni n-gram successor tables; an optional closed-class grammar skeleton), and **Tier B**, the attention+MLP composition kernel as plain matmuls. Faithfulness is enforced by a top-1 gate: Tier A reproduces the reference idiom predictions with zero per-position mismatches over 500 positions, and Tier B matches the numpy/torch reference exactly in f32 (e.g. GPT-2 200/200, Qwen2.5 32/32). All probes below are *explain-only*: the decode path is untouched, so there is no faithfulness-gate risk and the numbers describe the real model, not a surrogate.

3.2 Objects

Fix a model with hidden width d , vocabulary V , and unembedding rows $\{U_v \in \mathbb{R}^d\}_{v \in V}$. For a position, let $r \in \mathbb{R}^d$ be the final residual stream; the model’s logits are $L_v = \langle r, U_v \rangle + b_v$ and its prediction is $t = \arg \max_v L_v$. The residual is an additive sum of component writes $r = \sum_j d_j$ (embedding, each layer’s attention block, each layer’s MLP, down to individual heads and neurons), so each *source* j contributes $c_j^v = \langle d_j, U_v \rangle$ to every proposition and $L_v = \sum_j c_j^v$ exactly. The four measured quantities that organise everything are:

- the **participation ratio** $\text{PR} = (\sum_j c_j^t)^2 / \sum_j (c_j^t)^2$, the effective number of sources carrying the winning logit;
- the **normalised margin** $\Delta / \|U_t - U_{v^*}\|$ with $\Delta = L_t - L_{v^*}$ and v^* the runner-up, i.e. the Euclidean distance from r to the nearest power-diagram facet (Section 5.2);
- the **readout multiplicity** $\mu_t = \#\{\text{top-DLA sources } j :$

Model (store)	RETRIEVED	SELECTED	COMPOSED
Qwen-0.5B $\times 2$ (captured)	~ 25	~ 60	~ 15
Qwen-0.5B (corpus)	26.4	60.8	12.8
Pythia-70M (corpus)	35.6	56.6	7.8
Pythia-160M (corpus)	35.8	54.2	10.0
Pythia-410M (corpus)	36.8	51.6	11.6
Pythia-1B (corpus)	35.6	49.6	14.8

Table 1: Route split (% of next-token decisions). Retrieval dominates (RETRIEVED+SELECTED ≈ 87 –92%) across architectures and store constructions; the irreducible computation (COMPOSED) grows with scale against a fixed store.

$\arg \max_v c_j^v = t\}$, how many sources would, in isolation, already pick t ;

- the **differential incidence** (pivotality) $D_j = c_j^t - c_j^{v^*} = \langle d_j, U_t - U_{v^*} \rangle$, the amount by which ablating source j shifts the t -vs- v^* margin.

3.3 Models, holdouts, probes

The primary study uses **Qwen2.5-Coder-0.5B-Instruct** and **Qwen2.5-0.5B-Instruct** (shared vocabulary, hence a shared model-captured store), on a natural-text and a code holdout, ~ 300 –500 contexts at context window 64. The cross-architecture replication uses the **Pythia** ladder (70M/160M/410M/1B) under a corpus n-gram store held constant across architectures. The probes, each a `fieldrun` CLI mode, are: `--attribute` (the three-way route split), `--probe` (is selection a function of the firing state?), `--probe-dla` (PR, margin, μ_t), `--probe-facet` (exact nearest power-diagram facet over the full vocabulary), `--probe-ablate` (single- and multi-circuit causal ablation), `--probe-reconstruct` (per-block logit reconstruction), and `--probe-quant` (per-block quantisation sensitivity). Numbers are indicative ($n \sim 300$), not high-precision; the qualitative structure is what replicates.

4. Empirical Findings

4.1 The irreducible computation: a measured three-way split

Classifying each next-token decision by route gives, on natural text, roughly **25% RETRIEVED, 60% SELECTED, 15% COMPOSED**: the model’s labour is retrieval-dominated, and most of what looks like “computation” is disambiguation *within* a retrieved candidate set. The genuinely-computed irreducible computation is the $\sim 15\%$ (the COMPOSED route). The split is regime-dependent in the expected direction: at ~ 540 store candidates, the candidate set covers the model’s $\arg \max \sim 85\%$ of the time on natural text but only $\sim 63\%$ on code, because code is computed, not retrieved. Table 1 shows the split is robust across a second architecture family and a different store construction.

4.2 No selection in magnitude: a uniform 45-way sum

The obvious hypothesis, that retrieved tokens are decided by a dominant circuit and composed tokens are not, is false.




Figure 1: The same route split, visualised. Retrieval (RETRIEVED+SELECTED) dominates every model; the COMPOSED band (the irreducible computation, red, value annotated) grows monotonically up the Pythia ladder against a fixed store (7.8 $\rightarrow$ 14.8%), even as the geometry and margin structure stay scale-stable (Section 4).

The participation ratio is $PR \approx 42\text{--}49$ on Qwen and is *route-invariant*: RETRIEVED, SELECTED, and COMPOSED tokens all have the same ~ 45 -way distributed sum, even for tokens a single n -gram rule reproduces perfectly. No source dominates the logit magnitude, *ever*. The mechanism is a uniformly distributed additive sum followed by an argmax, exactly the readout superposition predicts [9]. On the Pythia ladder, PR runs $13 \rightarrow 27 \rightarrow 63 \rightarrow 36$ (70M/160M/410M/1B): not monotone in parameters, but monotone in head count (48/144/384/128). The “ ~ 45 -way sum” is thus an architecture-shape parameter (how many circuits exist to spread over), not a universal constant.

4.3 The two separating axes: margin and readout multiplicity

If magnitude does not separate the routes, two other quantities do, and they are nearly independent (per-position correlation $r \approx 0.15$, $\sim 2\%$ shared variance). The first is **geometry**: the exact nearest power-diagram facet distance is monotone $\text{RETRIEVED} \gg \text{SELECTED} > \text{COMPOSED}$ (coder 2.23/1.34/1.03, instruct 2.78/1.45/1.22; same shape on every Pythia rung), so retrieved tokens sit deep inside their decision cell and composed tokens sit near a facet. The second is **redundancy**: the readout multiplicity μ_t . Coverable tokens are *redundantly multiply-realised* (SELECTED mean $\mu_t = 1.45/1.06$), while $\sim 80\%$ of COMPOSED tokens are *strictly emergent*: $\mu_t = 0$, the answer is the argmax of the ~ 45 -way sum but of *no* summand (COMPOSED $\mu_t=0$ fraction 84%/76%). This is the operational definition of emergence we will formalise three ways in Section 5:

the argmax of a sum that is the argmax of no summand.

The split survives a confidence control: within matched margin bins, the covered-versus-composed redundancy gap persists (coder 65/71% vs 17/16%), so it is not merely that the store covers the model’s confident predictions.

A second, geometric *critical test* rules out the most natural story for composition. One might guess that a composed token is the residual r crossing the facet *out of the store’s predicted cell*. It is not: the nearest facet is the bisector with the store’s own prediction for only 14%/8% of COMPOSED tokens. For $\sim 85\%$ of composed tokens the store’s prediction is not even the nearest competitor, so composition is a *non-local* divergence from the rule, not a near-miss. The $\sim 14\%$ near-miss subclass is one identifiable thing: closed-class and morphological coin-flips ($a \perp the, will \perp is, ; \perp, -ful \perp -y$), the grammar regime where the store is strongest, not novel computation.

4.4 Causal ablation: composed is fragile, and the repair is diffuse

Turning the readout into an intervention, zeroing the top-DLA circuit in the forward pass and asking whether the prediction flips, converts these correlations into causal structure. Five results stand out.

(i) **Route-ordered fragility, margin-governed.** $\text{flip}@k=1$ is $\text{RETRIEVED } 22/26\% < \text{SELECTED } 40/48\% < \text{COMPOSED } 54/61\%$: knock out the single top circuit and a composed token flips $\sim 2.4\times$ as often as a retrieved one. But this tracks the margin, and a margin-matched split shows the margin is the governor: flips collapse $76 \rightarrow 41 \rightarrow 17\%$ across margin terciles *identically* for $\mu_t \geq 2$ and $\mu_t = 0$ tokens.

(ii) **Redundancy is non-compensatory.** Holding the ablated circuit fixed to a confirmed t -supporter and splitting on whether a backup remains ($\mu_t=1$ vs $\mu_t \geq 2$), backups confer no robust protection (pooled flip: low- Δ 90%/80%, high- Δ 4%/21%). Removing one supporter is not caught by the others: apparent agreement among many weak readers is *not* fault tolerance. The redundant supporters, individually $< 10\%$ of the logit at $PR \approx 40$, provide essentially no cushion.

(iii) **The causal variable is pivotality D_j versus margin Δ , not μ_t .** The *linear flip identity* $\text{flip} \iff \Delta < D_j$ holds as a near-perfect *necessary* condition: $\text{sign}(D_j - \Delta)$ mispredicts a non-flip only 3/300 (coder) and 17/300 (instruct) times. A logistic control flip $\sim \Delta + D_j + \mathbf{1}[\mu_t \geq 2]$ gives standardised weights $\Delta = -4.21/-3.11$, $D_j = +2.82/+1.16$, and μ_t essentially zero ($-0.60/+0.06$, opposite signs across models); dropping μ_t costs $+0.0035/+0.000$ mean log-loss. So μ_t is a *proxy* for (Δ, D_j) position, not an independent cause.

(iv) **Non-truth-functionality is a kernel.** With $\rho = \cos(U_t, U_{v^*})$ the runner-up coherence, pivotality and flip rate both fall as competitors become near-synonyms (coder $|D_j|$ $1.47 \rightarrow 0.86$, flip $53 \rightarrow 18\%$): common-mode evidence cancels

in $D_j = \langle d_j, U_t - U_{v^*} \rangle$. The competition geometry is read directly off the frame Gram, the structural fact Section 5.1 builds on.

(v) The repair is diffuse, not a lever. When the linear identity predicts a flip, an indirect downstream recomposition *rescues* the original token about half the time, and this cushion scales with the margin (rescue $14 \rightarrow 39 \rightarrow 50 \rightarrow 61\%$ across Δ). But the rescue is not localisable: 18–34% of rescues survive ablating *every* single downstream block, and the per-head un-rescue rate is order-1/PR: measured 4.0–4.1% against $1/\text{PR} = 2.5\text{--}2.8\%$, a $\sim 1.5\times$ constant. There is no surgical repair target; the repair is distributed for the same reason the readout is ($\text{PR} \approx 45$).

4.5 Per-block reconstruction: exact, block-sparse, circuit-dense

Decomposing the predicting logit into its 49 per-block residual writes (embedding + attention + MLP per layer) and checking $\sum_{\text{blocks}} = L_t$ gives *floating-point exact* reconstruction (mean error $5.9 \times 10^{-6} / 7.2 \times 10^{-6}$). Residual-stream additivity holds, so the static decomposition reconstructs every logit exactly, the empirical face of the soundness theorem of Section 5.3. The decision is *block-sparse* (effective supporting blocks $\approx 8\text{--}10$ of 49; top block to flip $\sigma \approx 1.1\text{--}1.6$) but *circuit-dense within a block* ($\text{PR} \approx 45$ over a block’s ~ 14 heads). The readable program is compact at block granularity and bottoms out there for composed tokens; below the block it is the dense computed remainder.

4.6 Scaling and a research-to-speed bridge

On the Pythia ladder the robust core (the three-way split, route-invariant PR, the RETRIEVED $\gg$ {SELECTED, COMPOSED} margin geometry, the refuted critical test, and margin-governed fragility) replicates at every scale. Two findings are scale-dependent. The irreducible computation *grows* with scale ($7.8 \rightarrow 14.8\%$ up the ladder against the same store; Figure 1): bigger models compute more. And μ_t -redundancy is something models *grow into*: at 70M/160M, $\mu_t \approx 0$ everywhere, while the full transition appears between 160M and 410M. The power-diagram geometry (margin) is scale-stable from 70M; the readout multiplicity is scale-emergent then stable, consistent with μ_t being a readout property, not the causal variable. Finally, quantisation sensitivity tracks the split: single-block int4 flip rates are RETRIEVED 3.6/2.4% $\ll$ COMPOSED 5.9/18.8%, so the retrievable tier survives aggressive quantisation while the irreducible computation is where quantisation error concentrates. This is the research result with a deployment consequence (quantise the retrievable tier hard, protect the computed core).

5. One Theory in Three Interpretations

The measurements above constrain a theory tightly. It must be *additive* (the logit is an exact sum, Section 4.5); *cardinality-inert* (the count μ_t has no causal weight, only D_j and Δ do, Sec-

tion 4.4); carry a *non-truth-functional coupling* that hardens with $\rho = \cos(U_t, U_{v^*})$; have a *weighted-threshold* decision beyond Horn clauses ($\mu_t=0$ for most composed tokens); and possess a second, *diffuse, non-localisable repair layer* bounded by $1/\text{PR}$. We now give a structure with all five properties, and observe that it is a single object viewed as logic, geometry, or computation, according to a temperature/semiring parameter.

5.1 Interpretation I, Logic ($T=1$): Projective Incidence Calculus

Objects. Sources S (circuits) are vectors d_j in an inner-product space $\mathcal{H}$; propositions V (tokens) are directions $U_v \in \mathcal{H}$. The *projective pairing* $c_j^v = \langle d_j, U_v \rangle$ is the direct logit attribution; aggregated evidence is the residual $r = \sum_j d_j$ with logits $L_v = \langle r, U_v \rangle = \sum_j c_j^v$. Two objects carry the theory. The *Gram kernel* $G_{vw} = \langle U_v, U_w \rangle$ is the proposition frame’s inner-product structure, and the *differential incidence* $D_j^{t,v} = \langle d_j, U_t - U_v \rangle$ is the atomic causal quantity. Where Bundy’s incidence of a proposition is a *set* $i(v) \subseteq I$, ours is the linear functional $\langle \cdot, U_v \rangle$, and the overlap of two propositions is not a set intersection but the kernel value G_{vw} . This is the central move: propositions share an inner-product space, so their incidences overlap *intrinsically*, and the structure that probability is not truth-functional over is carried by G .

Connectives. Combination is signed linear accumulation in log-space, $L_v = \sum_j c_j^v$, equivalently a *product of experts*, each source multiplying a proposition’s mass by $\exp(c_j^v)$ [13]. The decision is a weighted threshold: $\arg \max_v L_v$ is the Laguerre power diagram of $\{U_v\}$ with weights $\|U_v\|^2$, and the normalised margin is the facet distance (Section 5.2). The inferential analogue of incidence-calculus bounding is a *coalition bound*: from a known subset $S' \subseteq S$, the decision margin is bounded by the partial sum $\sum_{j \in S'} D_j$ plus a residual bound; the measured coalition additivity ($\text{sign} \sum_j D_j$ predicts joint-ablation flips at $\sim 75\text{--}83\%$) is its empirical face.

Theorems. The desiderata are recovered as consequences, not posited.

Theorem 1 (Cardinality-inertness). *Under the projective pairing, the decision depends only on $\{D_j, \Delta\}$ and is invariant to the count μ_t .*

This recovers the measured fact that μ_t carries no causal weight given (Δ, D_j) .

Theorem 2 (Non-truth-functionality budget). *Competition hardness between t and v is a monotone function of $\rho_{tv} = G_{tv}/\sqrt{G_{tt}G_{vv}}$; as $\rho \rightarrow 1$, the differential incidence $D_j = \langle d_j, U_t - U_v \rangle$ becomes common-mode and collapses. PIC reduces to classical Boolean incidence calculus exactly when G is diagonal.*

This recovers the measured ρ -boundary and identifies Bundy’s orthogonal case as the diagonal- G limit.

Theorem 3 (Weighted-threshold expressivity). *COMPOSED conclusions (those with no single sufficient sub-conjunction, $\mu_t=0$) are exactly the formulae not expressible in Horn / $\cap$ - $\cup$ incidence calculus but expressible with the weighted-threshold connective.*

Theorem 4 (Recovered probability). *Maintain a measure over propositions with uniform base M_0 ; let each source reweight every proposition’s mass by $\exp(c_j^v)$. Then $m(v) = M_0 \exp(\sum_j c_j^v) = M_0 \exp(L_v)$ and $m(v) / \sum_w m(w) = \exp(L_v) / Z$, exactly the model’s softmax. The Gibbs measure is recovered as a PIC incidence frequency, exactly and parameter-free.*

Theorem 4 is the direct descendant of Bundy’s “probability as proportion of worlds.” Crucially, G need not be injected as correlated noise (which would trade off exact recovery): it is *already* structural as the Gram of the proposition frame, so it structures the static logits (Theorem 4) and the competition geometry (Theorem 2) at once, with no variance trade-off.

The second layer. PIC has a monotone additive core (the map $r \mapsto (L_v)$ and the threshold) and a non-monotone fixpoint closure. Model the forward pass as an iterated operator T on the evidence vector; the additive core is its linearisation at the operating point, and the diffuse repair (the irreducible computation) is the higher-order closure T^∞ – (linearisation). It appears only under intervention (the off-diagonal Jacobian) and is provably diffuse in the high-PR regime:

Theorem 5 (Diffuseness). *Any causal property realised as $E = \sum_m e_m$ with equitable $e_m \sim E/\text{PR}$ has single-source influence $O(1/\text{PR})$; hence no bounded-size PIC formula localises it, and $P(\text{single-module intervention alters } E) \approx 1/\text{PR}$.*

This recovers the measured per-head un-rescue rate (Section 4.4(v)), order-of-magnitude confirmed at 4.0–4.1% versus $1/\text{PR} = 2.5\text{--}2.8\%$. The $\sim 1.5\times$ excess over $1/\text{PR}$ is consistent with a mild within-layer correlation of repair (a layer’s heads share rescue), the one place the equitable-contribution approximation is loosest. The clean one-line statement of the architecture: *the recovered probability lives in the static decomposition (additive, exact); the irreducible computation lives in the intervention response (non-additive, diffuse); the same model exhibits both as different layers.*

5.2 Interpretation II, Geometry ($T=0$): the Tropical Decision Surface

The decision is $\arg \max_v (\langle r, U_v \rangle + b_v)$, so the *max-logit*

$$M(r) = \bigoplus_v (b_v \otimes x^{U_v}) = \max_v (\langle r, U_v \rangle + b_v)$$

is a *tropical polynomial* in r over the $(\max, +)$ semiring ($a \oplus b = \max(a, b)$, $a \otimes b = a + b$): its monomial exponents are the unembedding rows U_v , its Newton polytope $\text{conv}\{U_v\}$ is the set of tokens that can ever win, and its tropical hypersurface (where the max is attained by ≥ 2 monomials) is the decision

boundary [25, 17]. Two of our measurements are this geometry, exactly:

Proposition 1 (Cells are a power diagram). *The linear regions of M are the Laguerre power diagram of $\{U_v\}$ with weights from $(b_v, \|U_v\|^2)$; the cell containing r is the predicted token.*

Proposition 2 (Margin is tropical distance). *The normalised margin $(L_t - L_{v^*}) / \|U_t - U_{v^*}\|$ is the exact Euclidean distance from r to the nearest facet of the tropical hypersurface.*

These are not conjectures: --probe-facet computes the exact nearest facet over all 151,936 tokens, and the measured monotone ordering RETRIEVED $\gg$ SELECTED $>$ COMPOSED (Section 4) is the facet-distance signature. Emergence translates sharply: decomposing $M(r) = \max_v \sum_j c_j^v$, a position is RETRIEVED when the winning cell is already selected by a single source’s monomial (a dominated vertex, $\mu_t \geq 1$) and COMPOSED when the winning region is *interior*: no single source’s monomial attains the max, the cell exists only in the sum ($\mu_t = 0$). This is the tropical reading of “argmax of a sum that is the argmax of no summand.”

The irreducible computation as a tropical-rank floor. A flat retrieval table (“context key $\rightarrow$ stored next-token logits”) is a tropical map whose monomials are exactly its stored keys, one tropical term per row, tropical rank bounded by the table size. Composition is precisely the decision regions requiring monomials *not* in the table: sums of stored keys that create new interior cells. Writing $\rho_{\text{trop}}(\text{core})$ for the tropical rank of the core’s decision map and $\rho_{\text{trop}}(\text{KB})$ for the best flat table at matched coverage [8], the *irreducible computation is exactly the unremovable part of $\rho_{\text{trop}}(\text{core}) - \rho_{\text{trop}}(\text{KB})$* : the decision cells no lookup table reproduces because they are composed (interior) monomials. The measured COMPOSED fraction is the empirical shadow of this gap. This also explains *why* a linear (SVD) rank misranks the core: its hardness is the number of *tropical* monomials (decision cells), which a Frobenius rank does not measure, the same reason a data-aware low-rank update beats plain SVD at matched rank in the decompiler’s entangled-core experiments.

5.3 Interpretation III, Computation: Semiring-Weighted Datalog

The core’s next-token computation is semiring accumulation over a finite relational domain, evaluated bottom-up: contributions sum along the residual stream, then one competitive aggregation picks the token. That is exactly Datalog evaluated under a *provenance semiring* [12, 14]. We export the model as a semiring-weighted Datalog program Π : propositions V with directions U_v and Gram G ; sources S with pairings c_j^v ; the retrievable fragment as compact *stratified* clauses (an induction head is a recursive clause $\text{next}(T) : \text{-match_prefix}(P), \text{follows}(P, T)$; an n -gram is a weighted fact; grammar is a unary constraint); the computed fragment as a flagged dense aggregate; and a provenance evaluator with a semiring/temperature parameter. Reading Π

as a Functional Aggregate Query, the decision is the aggregation $\oplus_v \otimes_j$ (factor).

Datalog rather than Prolog is a semantics decision: provenance semirings are a Datalog construct (Prolog has no semiring parameter, and cut / negation-as-failure / clause order break the semiring laws); forward residual accumulation matches Datalog’s bottom-up least-fixpoint; aggregation is native; and the program is decidable and statically analysable (PTIME data complexity, a unique least fixpoint): a mathematical object one can bound and verify, rather than a procedure. The retrievable rules are finite-domain over tokens and need none of Prolog’s extra power; the irreducible computation is dense arithmetic Prolog could not compact either.

Theorem 6 (Two-temperature soundness). *Read Π as a semiring FAQ, with sources combined by $\otimes$ along the residual derivation (so the accumulated value of v is $\otimes_j c_j^v = \sum_j c_j^v = L_v$) and competing propositions combined by $\oplus$. Then*

- under the log-semiring ($\oplus = \text{log-sum-exp}$, $\otimes = +$), $\llbracket \Pi \rrbracket$ is the sum-product: the aggregate is $\log \sum_v \exp(L_v) = \log Z$ and the per-proposition share is $\exp(L_v)/Z = P(v)$, the model’s softmax distribution (Interpretation I, $T=1$);
- under the tropical semiring ($\oplus = \max$, $\otimes = +$), $\llbracket \Pi \rrbracket$ is the max-product: the aggregate is $\max_v L_v$ with witness $\arg \max_v L_v$, the model’s greedy decode (Interpretation II, $T=0$).

This is the standard sum-product / max-product duality instantiated on Π , plus the product-of-experts recovery of Theorem 4; the $\log \leftrightarrow \max$ homomorphism is Maslov dequantization, so the export is correct at both temperatures by one identity. It is confirmed numerically by the exact per-block reconstruction (Section 4.5): $\sum_{\text{blocks}} = L_t$ to floating point means the static export is faithful. The emitter is implemented: `fieldrun export --logic` writes a runnable, Soufflé-compatible semiring-Datalog program for one decode (candidate facts, the Tier-A retrievable clauses, per-block contrib facts with the dense remainder folded into a "rest" block, and the $(\max, +)$ decode), self-checking $\sum_{\text{contrib}} = L_t$ and that the decode equals the model’s token; a built-in evaluator runs it under either semiring (`eval --semiring max` returns the token, `--semiring log` the distribution): *one program, two semirings, two temperatures*.

The provenance gap. The clean theorem assumes independent factors; they are not, because $G_{vw} = \langle U_v, U_w \rangle$ couples every pair. Hence:

Conjecture 1 (Provenance gap). *Faithful correlated evaluation of Π requires provenance valued in the frame geometry (carrying the U_v directions, or the operator G), not scalars in a commutative semiring over $\mathbb{R}$. Scalar semiring Datalog is exact only on diagonal G (the orthogonal / classical-incidence-calculus case); for dense G , scalar provenance drops the joint exactly as classical incidence calculus drops $i(A) \cap i(B)$.*

This is the same wall as the irreducible computation seen from the computation side: a dense G coupling is a high-treewidth factor graph, and sum-product is exponential in treewidth, so the correlated evaluation is at once intractable and non-compact. Solving the provenance gap would solve Interpretation I’s non-truth-functionality theorem; that one-to-one correspondence is the strongest evidence the three interpretations are one theory.

5.4 The unification

The three interpretations are one object under a single degree of freedom, the temperature, which is the semiring choice. As $T \rightarrow 0$, $T \log \sum_v \exp(L_v/T) \rightarrow \max_v L_v$ and $\text{softmax}(L/T) \rightarrow \arg \max$, by Maslov dequantization [16]. The power diagram is $\lim_{T \rightarrow 0}$ of the softmax cells; Interpretation I’s coherence kernel $\rho_{iv} = \cos(U_i, U_v^*)$ becomes the tropical facet angle; Interpretation I’s smoothed competition is the $T > 0$ viscosity regularisation of the tropical variety. The same irreducible computation is the high-PR residual (Interpretation I), the high tropical rank (Interpretation II), and the high-treewidth dense- G region (Interpretation III). *One wall, three measures: participation ratio, tropical rank, treewidth*. And the same kernel G is the non-truth-functionality operator, the facet geometry, and the provenance coupling. The retrievable fragment is, in all three, the additive / low-rank / low-treewidth part that exports to a compact verifiable program; the computed fragment is the diffuse / high-rank / dense- G remainder with no compact symbolic form. This is the program’s logic-programming thesis made precise: *the model is a semiring-weighted Datalog program, decoding is provenance evaluation, and its two temperatures are the softmax measure and the greedy argmax over one program*.

6. Reproducibility

Every quantitative claim traces to an explain-only probe in the `fieldrun` toolchain (the Apache-2.0 licensed source, bundles, and the exact holdout and store manifests behind Table 1 to be released with the camera-ready version), which runs a model from a flat bundle in pure Rust with no deep-learning framework. After converting a checkpoint (`fieldrun convert --model <hf-id> --arch rope`) and building a store, the probes reproduce the paper:

```
# three-way route split (Sec. 5.1)
fieldrun --bundle <m> --ids holdout.json \
  --store store.json --attribute
# PR, normalized margin, mu_t (Sec. 5.2-5.3)
fieldrun ... --probe-dla --n-eval 500
# exact nearest power-diagram facet (Sec. 6.2)
fieldrun ... --probe-facet
# causal ablation; mu_t-falsifier; 1/PR lemma
fieldrun ... --probe-ablate --n-eval 300 \
  --head-sweep
# per-block reconstruction: sum(blocks)==logit
fieldrun ... --probe-reconstruct --n-eval 300
# quantization sensitivity by route (Sec. 5.7)
fieldrun ... --probe-quant --bits 4 --n-eval 80
```



```
# emit a runnable semiring-Datalog program and
# evaluate it under either semiring (Sec. 6.3)
fieldrun ... export --logic --ctx 32 \
  --candidates 24 --out decode.dl
fieldrun eval decode.dl --semiring max # decode
fieldrun eval decode.dl --semiring log # softmax
```

The cross-architecture replication (Table 1) uses the same probes on Pythia bundles (`--arch neox`) under a corpus n-gram store, with the faithfulness gate validated top-1 against a pure-numpy reference. All probes leave the decode path untouched.

7. Open Problems

The theory is anchored where it can be measured and conjectured exactly where the measurements stop. The sharpest open questions are: **(O1)** the soundness and completeness of weighted incidence resolution, the coalition bound $\sum_{j \in S'} D_j$ as a sound inference rule; **(O2)** whether the support number $\sigma(i)$ (the smallest sufficient circuit set) scales as the participation ratio, tying cardinality-inertness to the absence of a small sufficient body, **(O3)** a derivation of the measured $\sim 1.5 \times$ constant in $P(\text{single-module repair}) = \kappa/\text{PR}$ from the fixpoint closure, accounting for the within-layer correlation of repair; **(O4, the central one)** a non-scalar, geometry-valued provenance semiring that carries G faithfully and reduces to the scalar log/tropical semiring on diagonal G (solving it closes the provenance gap and Interpretation I’s non-truth-functionality theorem at once); **(O5)** the treewidth of the core’s factor graph as a third quantitative irreducible-computation measure, related to the participation ratio and the tropical rank (one wall, three measures); and **(O6)** cross-scale invariance of G ’s spectrum and the retrievable/computed split as a program-wide thesis. A PIC *syntax* (compiling formulae directly from DLA traces, with composed positions flagged as “no compact formula”) would make the retrievable fragment a statically verifiable extracted program and the irreducible computation a named, bounded region.

8. Discussion and Conclusion

A transformer, asked for a token, spends most of its labour retrieving and selecting, and a small, scale-growing remainder computing something new. We have made that distinction measurable on real models, given it an exact geometric and causal characterisation, and shown that the same structure is a probabilistic logic, a tropical geometry, and an executable Datalog program according to a single degree of freedom, the temperature. The retrievable fragment is additive, low-rank, low-treewidth, and compactly verifiable; the irreducible computation is the diffuse, high-rank, dense-Gram region that no lookup table reproduces and no single module localises. The three interpretations are not analogies bolted together: they share one kernel G , one wall measured three ways, and one homomorphism (Maslov dequantization) between their temperatures.

The lineage is the point. Incidence calculus was built in 1985 on the insight that probability is not truth-functional but incidence is, and that one recovers probability by measuring incidences [4]. Forty years on, a transformer’s composition core turns out to be exactly that calculus with the orthogonality assumption removed: its propositions live in an inner-product space whose Gram kernel is the explicit carrier of the non-truth-functional joint, its connective is a weighted threshold rather than Boolean intersection, and its output probabilities are recovered, exactly, as incidence frequencies. The classical calculus sits inside ours as the diagonal- G limit. We have only added a substrate, a thermometer, and a fixpoint engine; the central idea is Bundy’s.

Acknowledgements

This work is, before anything else, a tip of the hat to Alan Bundy’s incidence calculus. It also owes a real debt to two AI collaborators: Grok, for brainstorming and the mathematical derivations (the product-of-experts recovery and the correlated-alternatives framing among them), and Claude (Opus and Fable), for the experimental design, implementation, and execution of the `fieldrun --probe-*` program from which the empirical constraints come. J. Allan Scott carried out this work independently and in a personal capacity; it was not sponsored by, and does not represent the views of, his employer, Vista Higher Learning (VHL).

References

- [1] Franz Aurenhammer. Power diagrams: Properties, algorithms and applications. *SIAM Journal on Computing*, 16(1):78–96, 1987.
- [2] Stella Biderman, Hailey Schoelkopf, Quentin Anthony, Herbie Bradley, Kyle O’Brien, Eric Hallahan, Mohammad Aflah Khan, Shivanshu Purohit, USVSN Sai Prashanth, Edward Raff, et al. Pythia: A suite for analyzing large language models across training and scaling. In *International Conference on Machine Learning (ICML)*, pages 2397–2430, 2023.
- [3] Trenton Bricken, Adly Templeton, Joshua Batson, Brian Chen, Adam Jermy, Tom Conerly, Nick Turner, Cem Anil, Carson Denison, Amanda Askell, et al. Towards monosemanticity: Decomposing language models with dictionary learning. *Transformer Circuits Thread*, 2023. <https://transformer-circuits.pub/2023/monosemantic-features/index.html>.
- [4] Alan Bundy. Incidence calculus: A mechanism for probabilistic reasoning. *Journal of Automated Reasoning*, 1(3):263–283, 1985.
- [5] Alan Bundy. The incidence calculus. *Encyclopedia of Artificial Intelligence*, 1992.
- [6] Stefano Ceri, Georg Gottlob, and Letizia Tanca. *Logic Programming and Databases*. Springer, 1990.

- [7] Hoagy Cunningham, Aidan Ewart, Logan Riggs, Robert Huben, and Lee Sharkey. Sparse autoencoders find highly interpretable features in language models. In *International Conference on Learning Representations (ICLR)*, 2024.
- [8] Mike Develin, Francisco Santos, and Bernd Sturmfels. On the rank of a tropical matrix. *Combinatorial and Computational Geometry*, 52:213–242, 2005.
- [9] Nelson Elhage, Tristan Hume, Catherine Olsson, Nicholas Schiefer, Tom Henighan, Shauna Kravec, Zac Hatfield-Dodds, Robert Lasenby, Dawn Drain, Carol Chen, et al. Toy models of superposition. *Transformer Circuits Thread*, 2022. https://transformer-circuits.pub/2022/toy_model/index.html.
- [10] Nelson Elhage, Neel Nanda, Catherine Olsson, Tom Henighan, Nicholas Joseph, Ben Mann, Amanda Askell, Yuntao Bai, Anna Chen, Tom Conerly, et al. A mathematical framework for transformer circuits. *Transformer Circuits Thread*, 2021. <https://transformer-circuits.pub/2021/framework/index.html>.
- [11] Mor Geva, Roei Schuster, Jonathan Berant, and Omer Levy. Transformer feed-forward layers are key-value memories. *Proceedings of EMNLP*, 2021.
- [12] Todd J. Green, Grigoris Karvounarakis, and Val Tannen. Provenance semirings. In *ACM Symposium on Principles of Database Systems (PODS)*, pages 31–40, 2007.
- [13] Geoffrey E. Hinton. Training products of experts by minimizing contrastive divergence. *Neural Computation*, 14(8):1771–1800, 2002.
- [14] Mahmoud Abo Khamis, Hung Q. Ngo, and Atri Rudra. FAQ: Questions asked frequently. In *ACM Symposium on Principles of Database Systems (PODS)*, pages 13–28, 2016.
- [15] Yann LeCun, Sumit Chopra, Raia Hadsell, Marc’Aurelio Ranzato, and Fu Jie Huang. A tutorial on energy-based learning. In *Predicting Structured Data*. MIT Press, 2006.
- [16] Grigory L. Litvinov. The maslov dequantization, idempotent and tropical mathematics: A brief introduction. *Journal of Mathematical Sciences*, 140(3):426–444, 2007.
- [17] Diane Maclagan and Bernd Sturmfels. *Introduction to Tropical Geometry*, volume 161 of *Graduate Studies in Mathematics*. American Mathematical Society, 2015.
- [18] Daniel McFadden. *Conditional Logit Analysis of Qualitative Choice Behavior*. Frontiers in Econometrics. Academic Press, 1974.
- [19] Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. Locating and editing factual associations in GPT. *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
- [20] Catherine Olsson, Nelson Elhage, Neel Nanda, Nicholas Joseph, Nova DasSarma, Tom Henighan, Ben Mann, Amanda Askell, Yuntao Bai, Anna Chen, et al. In-context learning and induction heads. *Transformer Circuits Thread*, 2022. <https://transformer-circuits.pub/2022/in-context-learning-and-induction-heads/index.html>.
- [21] Matthew Richardson and Pedro Domingos. Markov logic networks. *Machine Learning*, 62(1–2):107–136, 2006.
- [22] David S. Touretzky and Geoffrey E. Hinton. A distributed connectionist production system. *Cognitive Science*, 12(3):423–466, 1988.
- [23] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.
- [24] An Yang, Baosong Yang, Binyuan Hui, et al. Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*, 2024.
- [25] Liwen Zhang, Gregory Naitzat, and Lek-Heng Lim. Tropical geometry of deep neural networks. In *International Conference on Machine Learning (ICML)*, pages 5824–5832, 2018.